

Diffusion-Based Image Denoising

Generative AI – Mini Project

Group Members:

Saksham Wane(2411CS04)

Sourav Kumar (2411CS14)

Prantik Maji (2411CS16)

Arpit Kumar (2411CS23)

Problem Statement & Objective

- **Goal:** Apply a pretrained diffusion model (DDPM) to denoise synthetically noised images from CIFAR-10 and understand the image generation process across inference steps.
- **Motivation:**
 - Image denoising is crucial for low-light photography, medical imaging, and enhancing degraded visual data.
 - Diffusion models have emerged as powerful generative models capable of high-quality image synthesis.
- **Objectives:**
 - Add synthetic Gaussian noise to test images.
 - Iteratively denoise using a pretrained DDPM.
 - Visualize and evaluate the step-wise reconstruction.
 - Analyze output quality using metrics like PSNR and SSIM.

- **Denoising Diffusion Probabilistic Models (DDPMs):**

- Generative models that learn to reverse a noising process.
- Training involves learning to denoise step-wise from pure noise to a clean image.

- **Key Paper:**

- Ho et al., 2020 — "*Denoising Diffusion Probabilistic Models*".
- Introduced DDPMs as an alternative to GANs, offering stable training and high-quality outputs.

- **Core Equation:**

$$\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_{\theta}(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$$

- **Where:** ϵ_{θ} is the noise estimator network, typically a U-Net.

- **Dataset: CIFAR-10**

- 60,000 32x32 color images across 10 classes.
- We use the **test set** (10,000 images) for our experiments.
- Images were normalized and Gaussian noise ($\sigma = 0.3$) was added for denoising evaluation.

- **Tools and Libraries:**

- **PyTorch:** Core deep learning framework for loading and preprocessing CIFAR-10.
- **HuggingFace Diffusers:** Provided pretrained DDPM and scheduler for reverse diffusion.
- **scikit-image:** Used to compute PSNR and SSIM for evaluation.
- **Matplotlib:** Visualization of denoising steps and timelines.

- **Approach:**

- Used a pretrained DDPM (UNet2DModel from HuggingFace) trained on CIFAR-10.
- Synthetic Gaussian noise was added to CIFAR-10 test images.
- The pretrained DDPM was applied in reverse (denoising) mode over 1000 time steps.

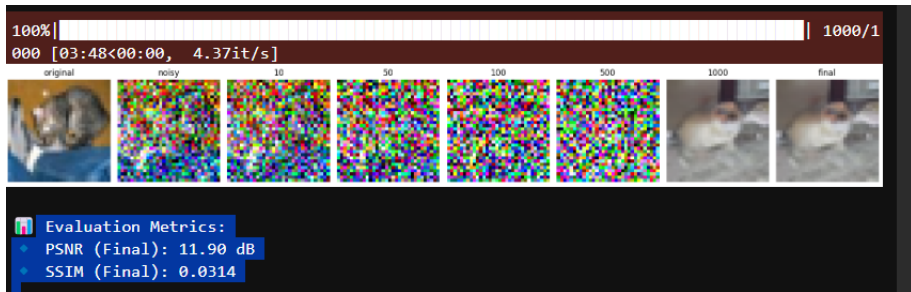
- **Implementation Pipeline:**

- 1 Load original CIFAR-10 test image and normalize.
- 2 Add Gaussian noise with standard deviation = 0.3.
- 3 Initialize DDPM scheduler and model.
- 4 Perform denoising over 1000 steps, saving snapshots at step 10, 50, 100, 500, and 1000.
- 5 Calculate final PSNR and SSIM.

- **Key Tools:** PyTorch, HuggingFace Diffusers, Matplotlib, scikit-image.

Experiments & Results

- **Visual Timeline:** Captured denoised image at intervals to show progression.
- **Selected Steps:** Step 10, 50, 100, 500, 1000, and final output.
- **Quantitative Evaluation:**
 - PSNR (Final): **11.90 dB**
 - SSIM (Final): **0.0314**
- **Findings:**
 - Minor structural improvements are noticeable over time.
 - Noise reduction is visible, but finer details are not fully restored.



- **Progressive Restoration:**

- The denoising process clearly reduces noise, especially after 100 steps.
- Major improvements observed around step 500, indicating non-linear convergence.

- **Quality Metrics Interpretation:**

- Low PSNR suggests limited pixel-wise accuracy.
- Very low SSIM reveals poor structural similarity to the original image.

- **Possible Causes:**

- Pretrained DDPM was not fine-tuned for higher noise levels or specific image details.
- Lack of conditioning (e.g., class labels or guidance) affects denoising specificity.

- **Conclusion from Analysis:**

- DDPM shows strong potential but requires contextual guidance or domain-specific tuning for high-fidelity results.

Challenges & Limitations

- **Model Limitations:**

- The pretrained DDPM is not optimized for arbitrary noise levels.
- Absence of class-conditioning limits recovery of semantic structures.

- **Performance Bottlenecks:**

- Full 1000-step inference is computationally expensive.
- Real-time or interactive applications are not feasible in current setup.

- **Evaluation Constraints:**

- PSNR and SSIM do not capture perceptual quality fully.
- Metrics may undervalue visually decent reconstructions with minor spatial shifts.

- **Improvements for Future:**

- Explore classifier-guided or diffusion transformers.
- Apply training-time augmentation or fine-tuning on noisy image sets.

Conclusion & Takeaways

- **Summary:**

- Implemented DDPM-based image denoising using a pretrained model.
- Visual and quantitative evaluation showed noise reduction but limited restoration fidelity.

- **What We Learned:**

- Diffusion models effectively smooth and denoise images over multiple iterations.
- Model architecture and conditioning play a critical role in output quality.

- **Future Scope:**

- Experiment with conditional DDPM or control-net guidance.
- Use real-world noisy datasets for evaluation and training.
- Compare against GAN-based or hybrid denoising pipelines.

- Ho et al. (2020). *Denoising Diffusion Probabilistic Models*.
- CIFAR-10: <https://www.cs.toronto.edu/~kriz/cifar.html>
- Diffusers Library: <https://huggingface.co/docs/diffusers>
- Code Repo
<https://github.com/VaruosKumar/GenAI-Mini-Project>